

Practical Work 12 – 16/05/2024

Non sequential architectures - Transfer learning - Autoencoders

Objectives

The first objective of this PW is to practice with non-sequential deep architectures. The second objective is to experiment with transfer learning. Optionnally, we also propose you with an exercise on autoencoders. To provide your solutions, you can use either the *functional API* of Keras (as seen in class) or build an equivalent using Pytorch.

Submission

- **Deadline** : Wednesday 24th of May, 12 :00 (noon)
- **Format** : Zip with report and iPython notebook.

Exercise 1 Non-sequential architectures

Using the CIFAR10 dataset :

- Re-use one of your best CNN architecture from PW04 or PW06 and transform the model definition to use non-sequential strategies such as *multiple features* or *multiple paths* as described in the slides.
- Optional : generate graphs of the architectures as shown in class. For that, you may install GraphViz to use the Keras function `plot_model()`, or an equivalent (note – you may need to launch the installation of pydot first with `pip install pydot`).
- Use *callbacks* to save the best trained models according to a monitoring of the accuracy on the test set.

Report on your experiments and describe your best configuration through experimenting with 2-3 different architectures. Use a table similar to the example given below and provide the hyper-parameters used for your configurations.

Model	Architecture	Callback	Acc. train %	Acc. test %
1	Description of architecture...	yes	84.2%	78.4%
2
3
4
5

TABLE 1 – Performances on CIFAR10 with different sequential and non-sequential configurations. In the Architecture column, specify the number of filters, their size, stride and padding values, etc.

Exercise 2 Transfer learning

With Keras and Pytorch, you have access to many pre-trained architectures on ImageNet. Such architectures are able to extract robust image features and can generally be used for many different image tasks. The objective of the exercise is to check if such features lead to good results on the Food101 dataset. This dataset contains 101'000 color images of meals corresponding to 101 classes. In order not to have too long trainings, only the first 20 classes will be used here.

For the one using Keras, the following tutorial can probably help you with the exercise : https://keras.io/guides/transfer_learning/#build-a-model

- Review the slides on using transfer learning.
- Chose one of the architecture presented here : <https://keras.io/applications/>. Beware that some architectures are using large memory and lots of cpu (so better move to gpu if you chose large models)!
- Download the provided notebook `ex2_transfer_learning_stud.ipynb`.
- Implement the preprocessing and the training phases.
- Experiment with one or several architectures with different classification heads (hidden-layers, ...).
- Describe the results in your report which architecture and parameters works best.

Exercise 3 Optional - Auto-encoders

We will use the notMNIST data set. First get the .gz files from this site : <https://github.com/davidflanagan/notMNIST-to-MNIST>.

Use as a starting point the notebook `notMNIST-auto-encoder-stud.ipynb`.

Create a shallow dense autoencoder

Using the Keras library, build a very simple autoencoder with one input layer, one output layer and one hidden layer.

- a) Chose the dimension of your hidden layer. As this is the size of your encoding, you should chose a “reasonable” dimension, in between the number of classes of your problem and the dimension of input and output layers¹.
- b) Define your layers, don’t forget that your input dimension is the same as your output.
- c) Create your models (using the keras `Model`) : encoder, decoder and autoencoder (encoder and decoder).
- d) Don’t forget to compile your autoencoder. Fit it on the data!
- e) Visualize your encoding image with the input images on the top and their corresponding encoded ones on the bottom. Your visualization may differ slightly from the one on Figure 1
- f) How could we evaluate the performance of such a “compression” tool (in terms of loss and in terms of gain of bandwidth)? You just need to comment – no need to compute anything here².
- g) Are you now able to use the encoded features to train a simple Support Vector Machine – SVM classifier? Comment on the performance by comparing the extracted features of the auto-encoder with using the pixel values as input of the SVM.




FIGURE 1 – Input and output of the autoencoder

Create a convolutional de-noising autoencoder

Create a model in Keras with the following structure :

- Encoder
 - Convolution layer with 32 filters of size 3 x 3
 - Downsampling (max-pooling) layer of size 2 x 2
 - Convolution layer with 64 filters of size 3 x 3
 - Downsampling (max-pooling) layer of size 2 x 2
 - Convolution layer with 128 filters of size 3 x 3
- Decoder
 - Convolution layer with 128 filters of size 3 x 3
 - Upsampling layer of size 2 x 2 (see `UpSampling2D()`)
 - Convolution layer with 64 filters of size 3 x 3
 - Upsampling layer of size 2 x 2

1. We tried with 32 and it was working well, you may try yourself with different values to observe the impact.
 2. If you really want to compute something, you can of course.

— Convolution layer with 1 filter of size 3 x 3

- a) Check with the `summary()` method that you have indeed a diabolio network and that your output shape is the same as your input shape.
- b) Would it make sense to use the output of the encoder to compress the images? Comment your answer.
- c) Generate noisy versions of your train and test data using (see examples in slides).
- d) Train your network using a MSE loss and batch sizes of 128.
- e) Visualise the denoising capacity on some test data.

Exercise 4 Optional - Review Questions

- a) Why (or when) do we need the functional API of Keras?
- b) What are the benefits of using transfer learning?
- c) What does *freezing* mean in a transfer learning strategy? What is the influence of the training data set size on the *freezing*?
- d) What are the potential usage of auto-encoders?
- e) In which situations would you use auto-encoders to extract features?
- f) What are the principles of self-supervised learning?
- g) In which situations would you prefer to use auto-encoder or self-supervised learning instead of using pre-trained networks to extract features?